

Computing the cosine using Taylor series approximations

Written by Zamfir Ștefan-Alexandru, 17 March 2022.

This project aims to quantify the precision of my implementation of the cosine function.

All results will be compared against MATLAB's implementation, $\cos(x)$.

Table of Contents

The <code>taylorcos</code> function.....	1
<code>taylorcos.m</code>	1
How the method compares to MATLAB's by number of Taylor terms.....	2
Limitations.....	4

The `taylorcos` function

There are 4 important optimizations I discovered while working on the function:

1. **The cosine is an even function.** All negative values of x can be converted to positive values using $\text{abs}(x)$.
2. **The cosine is a periodic function.** We only need the result of $\text{mod}(x, \pi * 2)$ in our computations, also saving a huge amount of time and precision.
3. The cosine in $[0, \pi * 2)$ is **composed as an y-reflection** of the cosine in $[0, \pi)$ and $[\pi, \pi * 2)$. If $x > \pi$, then we can substitute it with $\pi * 2 - x$.
4. The cosine in $[0, \pi)$ is **composed as a x-reflection** of the cosine in $[0, \pi / 2)$ and $[\pi / 2, \pi)$. if $x > \pi / 2$, then we can substitute it with $\pi - x$ **and multiply the end result by -1**.
5. If, after applying the previous steps, the x is in $[\pi / 4, \pi / 2]$, then we reduce it further by computing the sine. The appropriate formula is $\cos(x) = \sin(\pi / 2 - x)$. We use the Taylor series for the sine instead of cosine in order to compute the sine.

The reason for reducing the range of x to $[0, \pi / 4]$ is simple. The more terms we compute in the Taylor series, the bigger the approximation range gets. What I mean by that is you can't compute the cosine of $10 * \pi$ if you only have the first three terms of the series. You need to bring more terms for that.

Well, because our range is so small, we will get **an accurate approximation within the first 8 - 9 terms**. If we were to not use the 5th step, the approximation would only get good after the 10th term.

`taylorcos.m`

```
% Input is x and terms, these are example values
x = 123456.123456;
terms = 12;

% The cos function has a periodicity of 2 * pi
% Also, the cos function is an even function
x0 = abs(mod(x, 2 * pi));
coef = 1;

% First reflection
if x0 > pi
    x0 = 2 * pi - x0;
```

```

end

% Second reflection
if x0 > pi / 2
    coef = -1;
    x0 = pi - x0;
end

y = 0;
% Last optimization
if x0 > pi / 4
    x0 = pi / 2 - x0;

    % Sine
    for i = 0:terms
        y = y + (-1) ^ i * x0 ^ (2 * i + 1) / factorial(2 * i + 1);
    end
else
    % Cosine
    for i = 0:terms
        y = y + (-1) ^ i * x0 ^ (2 * i) / factorial(2 * i);
    end
end

% In case we reflected the second time
y = coef * y;

fprintf("taylorcos(x) using %g terms (above) and cos(x) (below)\n%.16d\n%.16d\n", ...
        terms, y, cos(x));

```

```

taylorcos(x) using 12 terms (above) and cos(x) (below)
-5.7601758031359596e-01
-5.7601758031752992e-01

```

```

fprintf("Their difference is:\n%.4d", abs(y - cos(x)));

```

```

Their difference is:
3.9340e-12

```

How the method compares to MATLAB's by number of Taylor terms

Let us plot the difference between the two methods:

```

terms = 1;
maxterm = 12; % The amount of terms computed by the Taylor series
maxdiffs = zeros(1, maxterm);

x = -7:.02:7; % Range of computation

% Iterate through the number of terms to be used by series
while terms < maxterm + 1
    outputs = zeros(1, 10/.02); % Output of taylorcos
    tests = zeros(1, 10/.02); % Output of cos
    diff = zeros(1, 10/.02); % Their difference

```

```

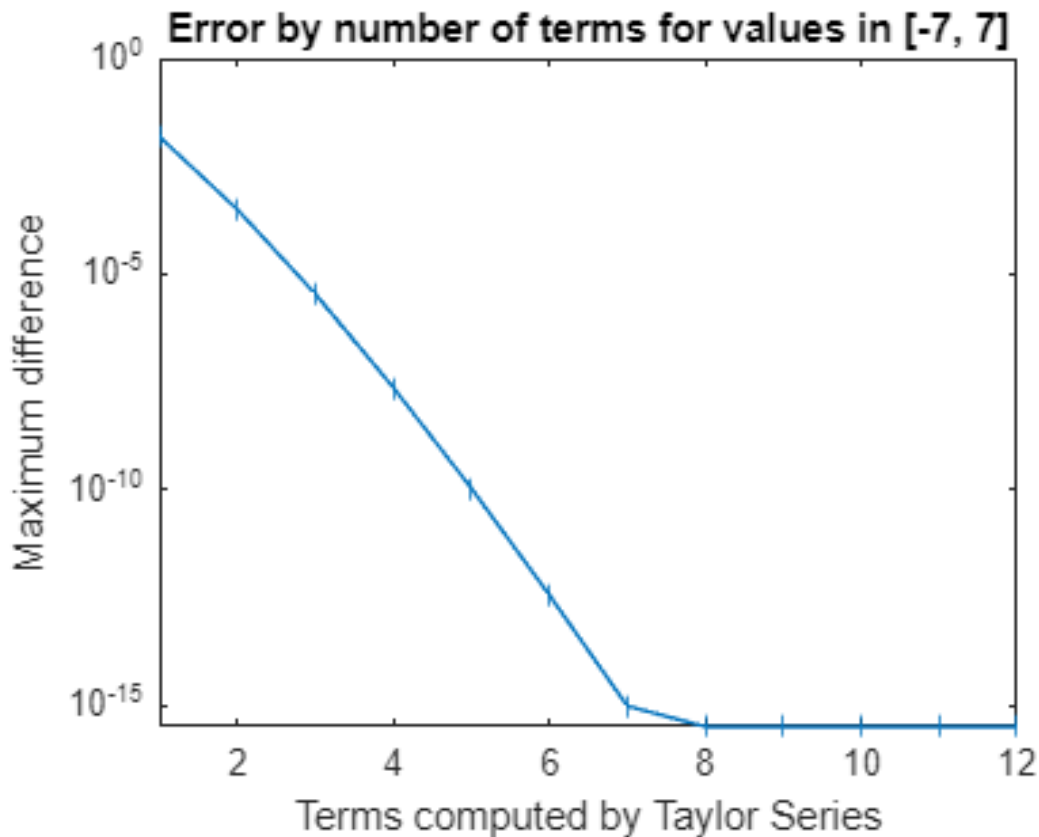
counter = 1;
for i = x
    outputs(counter) = taylorcos(i, terms);
    tests(counter) = cos(i);
    diff(counter) = outputs(counter) - tests(counter);

    counter = counter + 1;
end

maxdiffs(terms) = max(abs(diff));
terms = terms + 1;
end

% Plot the error by number of terms
plot(1:maxterm, maxdiffs, "-|");
title('Error by number of terms for values in [-7, 7]')
xlabel('Terms computed by Taylor Series')
ylabel('Maximum difference')
set(gca, 'YScale', 'log')
xlim([1 maxterm])

```



Interestingly, for the 10th, 11th and the 12th terms, the difference is exactly the same ($3.33\text{e-}16$). I suspect there simply isn't room for double to quantify numbers less than $3.33\text{e-}16$.

Limitations

The difference starts to broaden as the value of x increases since we are computing $\text{mod}(x, \pi * 2)$. This happens because representable floating point numbers start getting sparser and sparser as their magnitude increases. If $x > 10^{17}$, the modulo function starts to fail.

```
in = 10^17;
epsilon = eps(in);
fprintf("The difference between %d and the next represented double is %d", in, epsilon);
```

The difference between 10000000000000000 and the next represented double is 16

A difference of 16. You really can't provide the modulo of 6.38 if your numbers are 16 units apart.

The solution is to accurately compute the modulo. If that is done correctly, then the Taylor approximations are guaranteed to give you a difference no more than $3.33\text{e-}16$ from MATLAB's implementation. For your information, $\text{eps}(\pi / 4) = 1.11\text{e-}16$. The function is really close to MATLAB's.

There exists a way of doing that by providing enough bits of π . An implementation that's used by MATLAB can be found [here](#).